


sub14 — iso-architecture analysis

Loss vs compute, per model size

each curve = one architecture; ● = a run (LR-decayed)



Best loss vs architecture size

diminishing returns of width/depth

